

The POS tagger produces mostly correct tags, but a few of the tags are inaccurate (for example, in sentence 10150, “coming” is labeled as a noun when it should be a verb). The tagger is mostly accurate because it is trained on data that is similar to the testing data (they are from the same corpus) so we would expect the behavior / usage of most words to be the same in the testing and training data. However, there are still some discrepancies in usage, and we also have unknown words which introduce additional uncertainty into the labeling. Additionally, our tagger assumes that the POS at any given point in the sequence is affected only by the POS directly before it, and also that the word / observation at any given point is affected only by its POS and not by the preceding word(s) or POS(s), which is obviously not the case in English.